
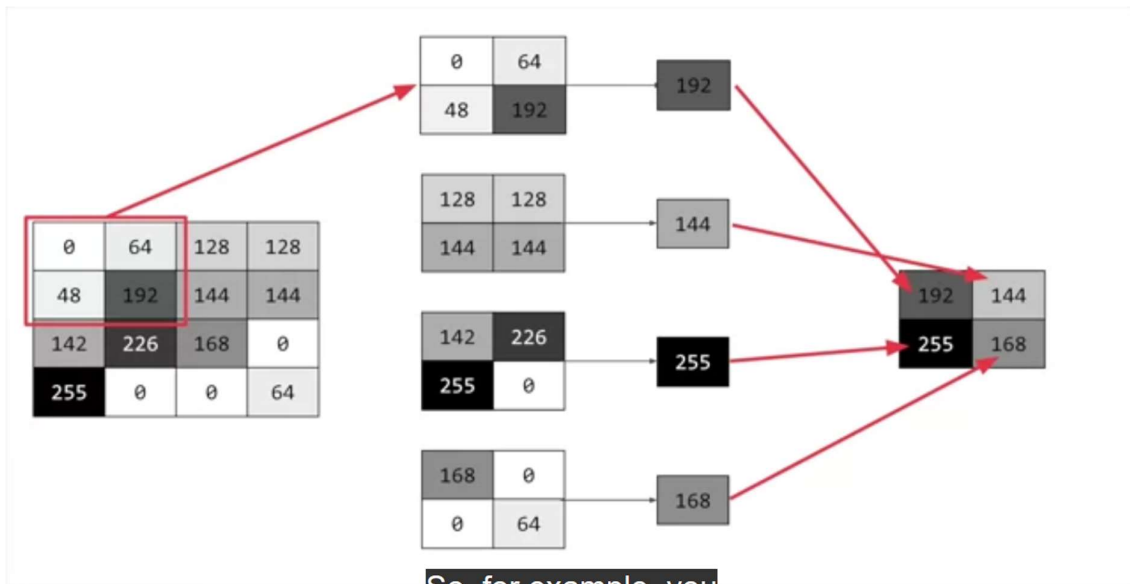


each corresponding filter value,



then the vertical lines in the image really pop out.



So, for example, you  
can see it here.

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation=tf.nn.relu),
    tf.keras.layers.Dense(10, activation=tf.nn.softmax)
])
```

So here's our code from the earlier  
example, where we defined out a neural

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(64, (3,3), activation='relu',
                           input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```


To add convolutions to this,  
you use code like this.

```
model.summary()
```

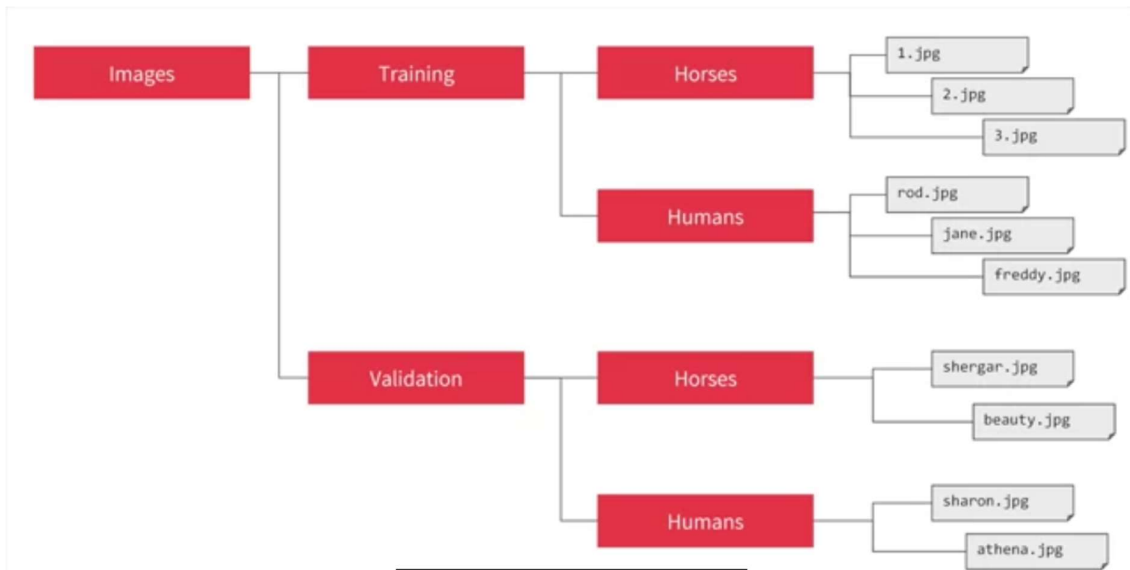
the model is the  
model.summary method.

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_12 (MaxPooling)	(None, 13, 13, 64)	0
conv2d_13 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_13 (MaxPooling)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_10 (Dense)	(None, 128)	204928
dense_11 (Dense)	(None, 10)	1290

It's important to keep an eye on the output shape column.



For example, how about these images of horses and humans?



So for example, consider  
this directory structure.

```
from tensorflow.keras.preprocessing.image  
import ImageDataGenerator
```

is available in  
Keras.preprocessing.image.

```
test_datagen = ImageDataGenerator(rescale=1./255)

validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(300, 300),
    batch_size=32,
    class_mode='binary')
```

same except of course it points  
at a different directory,

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3,3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2,2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

As you can see, it's the  
sequential as before with

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d_5 (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_6 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_6 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_7 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_7 (MaxPooling2D)	(None, 35, 35, 64)	0
flatten_1 (Flatten)	(None, 78400)	0
dense_2 (Dense)	(None, 512)	40141312
dense_3 (Dense)	(None, 1)	513

Total params: 40,165,409  
 Trainable params: 40,165,409  
 Non-trainable params: 0

Now, if we take a look  
at our model summary,

```

from tensorflow.keras.optimizers import RMSprop

model.compile(loss='binary_crossentropy',
              optimizer=RMSprop(lr=0.001),
              metrics=['acc'])

```

Now, you could do that again, but



```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=8,  
    epochs=15,  
    validation_data=validation_generator,  
    validation_steps=8,  
    verbose=2)
```

this looks a little different than  
before when you called model.fit.

```
import numpy as np  
from google.colab import files  
from keras.preprocessing import image  
  
uploaded = files.upload()  
  
for fn in uploaded.keys():  
  
    # predicting images  
    path = '/content/' + fn  
    img = image.load_img(path, target_size=(300, 300))  
    x = image.img_to_array(img)  
    x = np.expand_dims(x, axis=0)  
  
    images = np.vstack([x])  
    classes = model.predict(images, batch_size=10)  
    print(classes[0])  
    if classes[0]>0.5:  
        print(fn + " is a human")  
    else:  
        print(fn + " is a horse")
```

So these parts are specific to Colab,  
they are what gives